

Homework: User-level threads

In this assignment you will complete a simple user-level thread package by implementing the code to perform context switching between threads. Submit your solutions before the beginning of the next lecture to [the submission web site](#).

Switching threads

Download [uthread.c](#) and [uthread_switch.S](#) into your xv6 directory. Make sure `uthread_switch.S` ends with `.S`, not `.s`. Add the following rule to the xv6 Makefile after the `_forktest` rule:

```
_uthread: uthread.o uthread_switch.o
$(LD) $(LDFLAGS) -N -e main -Text 0 -o _uthread uthread.o uthread_switch.o $(ULIB)
$(OBJDUMP) -S _uthread > uthread.asm
```

Make sure that the blank space at the start of each line is a tab, not spaces.

Add `_uthread` in the Makefile to the list of user programs defined by `UPROGS`.

Run xv6, then run `uthread` from the xv6 shell. The xv6 kernel will print an error message about `uthread` encountering a page fault.

Your job is to complete `thread_switch.S`, so that you see output similar to this (make sure to run with `CPUS=1`):

```
~/classes/6828/xv6$ make CPUS=1 qemu-nox
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB, 4.9 MiB) copied, 0.0190287 s, 269 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes copied, 7.2168e-05 s, 7.1 MB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
291+1 records in
291+1 records out
149040 bytes (149 kB, 146 KiB) copied, 0.000528827 s, 282 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 1 -m 512
xv6...
cpu0: starting
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ uthread
my thread running
my thread 0x2D68
my thread running
my thread 0x4D70
my thread 0x2D68
my thread 0x4D70
my thread 0x2D68
my thread 0x4D70
my thread 0x2D68
...
```

`uthread` creates two threads and switches back and forth between them. Each thread prints "my thread ..." and then yields to give the other thread a chance to run.

To observe the above output, you need to complete `thread_switch.S`, but before jumping into `uthread_switch.S`, first understand how `uthread.c` uses `thread_switch`. `uthread.c` has two global variables `current_thread` and `next_thread`. Each is a pointer to a thread structure. The thread structure has a stack for a thread and a saved stack pointer (`sp`, which points into the thread's stack). The job of `uthread_switch` is to save the current thread state into the structure pointed to by `current_thread`, restore `next_thread`'s state, and make `current_thread` point to where `next_thread` was pointing to, so that when `uthread_switch` returns `next_thread` is running and is the `current_thread`.

You should study `thread_create`, which sets up the initial stack for a new thread. It provides hints about what `thread_switch` should do. The intent is that `thread_switch` use the assembly instructions `popal` and `pushal` to restore and save all eight x86 registers. Note that `thread_create` simulates eight pushed registers (32 bytes) on a new thread's stack.

To write the assembly in `thread_switch`, you need to know how the C compiler lays out `struct thread` in memory, which is as follows:

```
-----
| 4 bytes for state |
-----
| stack size bytes |
| for stack        |
-----
| 4 bytes for sp    |
----- <--- current_thread
.....
.....
-----
| 4 bytes for state |
-----
| stack size bytes |
| for stack        |
-----
| 4 bytes for sp    |
----- <--- next_thread
```

The variables `next_thread` and `current_thread` each contain the address of a struct `thread`.

To write the `sp` field of the struct that `current_thread` points to, you should write assembly like this:

```
movl current_thread, %eax
movl %esp, (%eax)
```

This saves `%esp` in `current_thread->sp`. This works because `sp` is at offset 0 in the struct. You can study the assembly the compiler generates for `uthread.c` by looking at `uthread.asm`.

To test your code it might be helpful to single step through your `thread_switch` using `gdb`. You can get started in this way:

```
(gdb) symbol-file _uthread
Load new symbol table from "/Users/kaashoek/classes/6828/xv6/_uthread"? (y or n) y
Reading symbols from /Users/kaashoek/classes/6828/xv6/_uthread...done.
(gdb) b thread_switch
Breakpoint 1 at 0x204: file uthread_switch.S, line 9.
(gdb)
```

The breakpoint may (or may not) be triggered before you even run `uthread`. How could that happen?

Once your `xv6` shell runs, type "`uthread`", and `gdb` will break at `thread_switch`. Now you can type commands like the following to inspect the state of `uthread`:

```
(gdb) p/x next_thread->sp
$4 = 0x4ae8
(gdb) x/9x next_thread->sp
0x4ae8 : 0x00000000 0x00000000 0x00000000 0x00000000
0x4af8 : 0x00000000 0x00000000 0x00000000 0x00000000
0x4b08 : 0x000000d8
```

What address is `0xd8`, which sits on the top of the stack of `next_thread`?

Submit: your modified `uthread_switch.S`

Optional challenges

The user-level thread package interacts badly with the operating system in several ways. For example, if one user-level thread blocks in a system call, another user-level thread won't run, because the user-level threads scheduler doesn't know that one of its threads has been descheduled by the `xv6` scheduler. As another example, two user-level threads will not run concurrently on different cores, because the `xv6` scheduler isn't aware that there are multiple threads that could run in parallel. Note that if two user-level threads were to run truly in parallel, this implementation won't work because of several races (e.g., two threads on different processors could call `thread_schedule` concurrently, select the same runnable thread, and both run it on different processors.)

There are several ways of addressing these problems. One is using [scheduler activations](#) and another is to use one kernel thread per user-level thread (as Linux kernels do). Implement one of these ways in `xv6`.

Add locks, condition variables, barriers, etc. to your thread package.